

Vitess

Getting Started with Vitess using Kubernetes

Dan Kozlowski

Minister Plenipotentiary of Engineering, PlanetScale

Database on Ephemeral Containers?

- Storage on Ephemeral nodes?
- Availability when resources can vanish?
- Configuration and service discovery



Build a cloud Native Database

- Should scale with demand
- Should be handled like row crops
- Should be logically singular
- Should be resilient in the face of failure
- Automated management



What does it take to run On K8s

- databases
- vttablet's
- vtgate's
- lockserver



The Global Control Plane

- Global Lock server
- vtctld - Controller Daemon and User Interface
- vtctlclient - Command line client to run commands



[demo]



Cluster Architecture

- Keyspace
 - Shards
 - Master
 - Cell
 - Replica
 - [Cell]
- Cell
 - Keyspace
 - Shard
 - Replica



The First Keyspace

- Register the Keyspace
 - CreateKeyspace
- Register the Cell
 - AddCellInfo



vtablet

- Tablets serve a shard
- They have a specific type
 - Ronly, replica, master
- They live in a cell
- They have a database backing them



Replication in Vitess

- Vitess uses the MySQL binlog
- vttablet consumes the binlog from other vttablets
- vitess uses this hijacking of the binlog to perform a number of tasks related to cluster management



[demo]



vtgate

- Presents a mysql interface (3306)
- Presents a custom grpc interface
- Can Talk cross shard
- Handles several activities related to sharding
 - Routes queries
 - Is responsible for some retries
 - Many other things

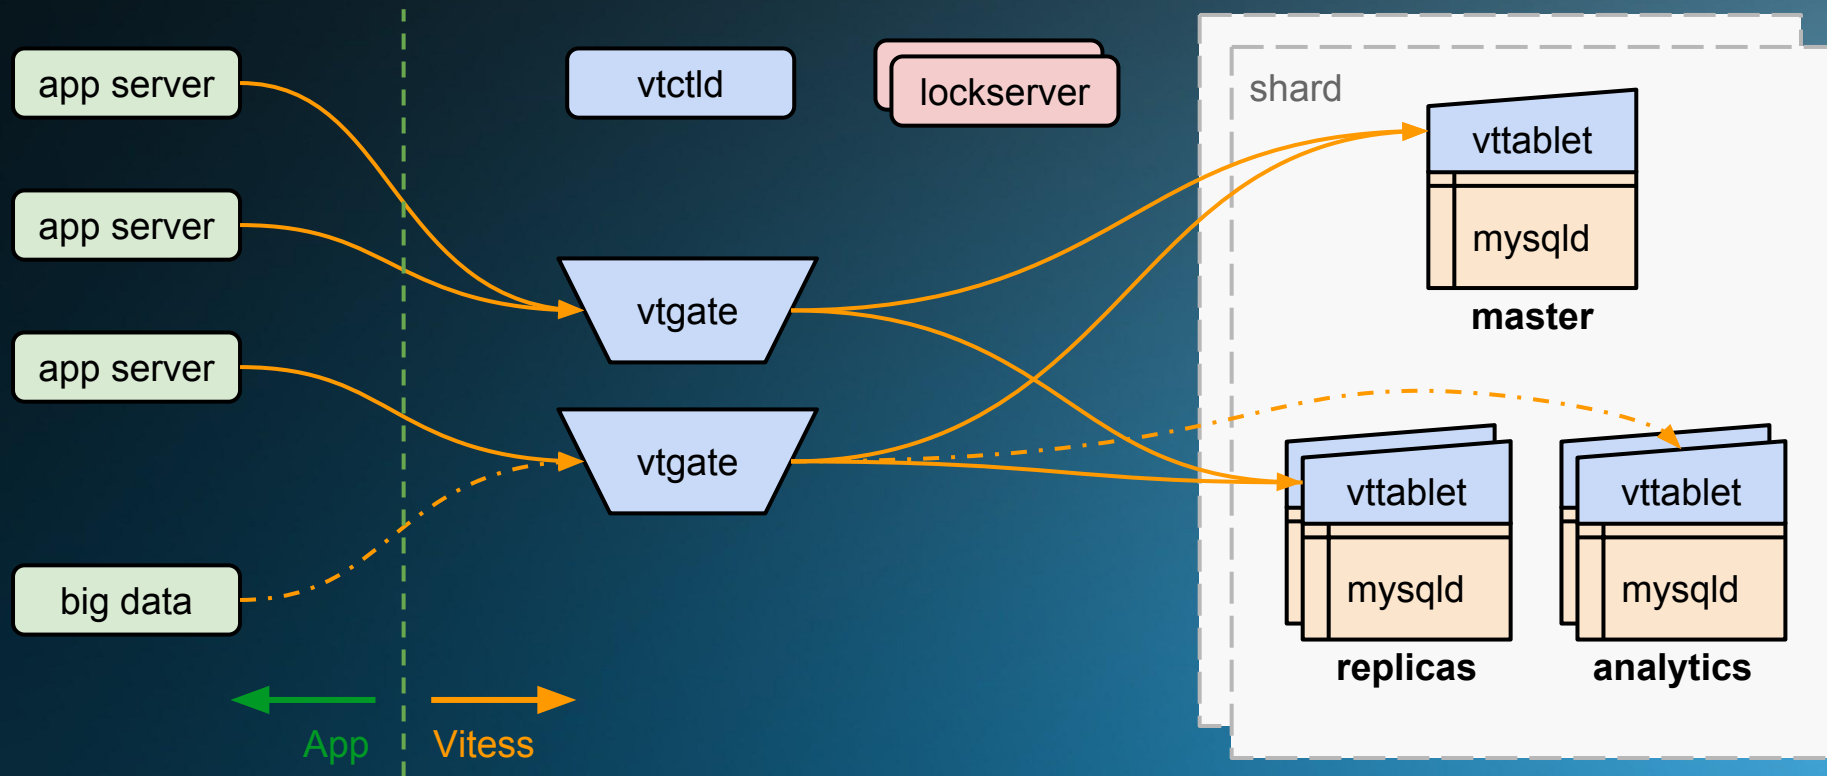


vtworker, Resharding, Schema Rollout

- Major workflow item is resharding
- All workflows should be done live
 - Vtgate knows how to handle the synchronous actions



Architecture



<https://github.com/planetscale/k8s-demo>



PlanetScale

contact@planetscale.com

SaaS / Support / k8s Operator



VReplication

It's all coming together!

Sugu Sougoumarane
Cofounder & CTO, PlanetScale

Vitess Resharding

- Split Copy
- Filtered Replication
- Failover
- Well used and robust

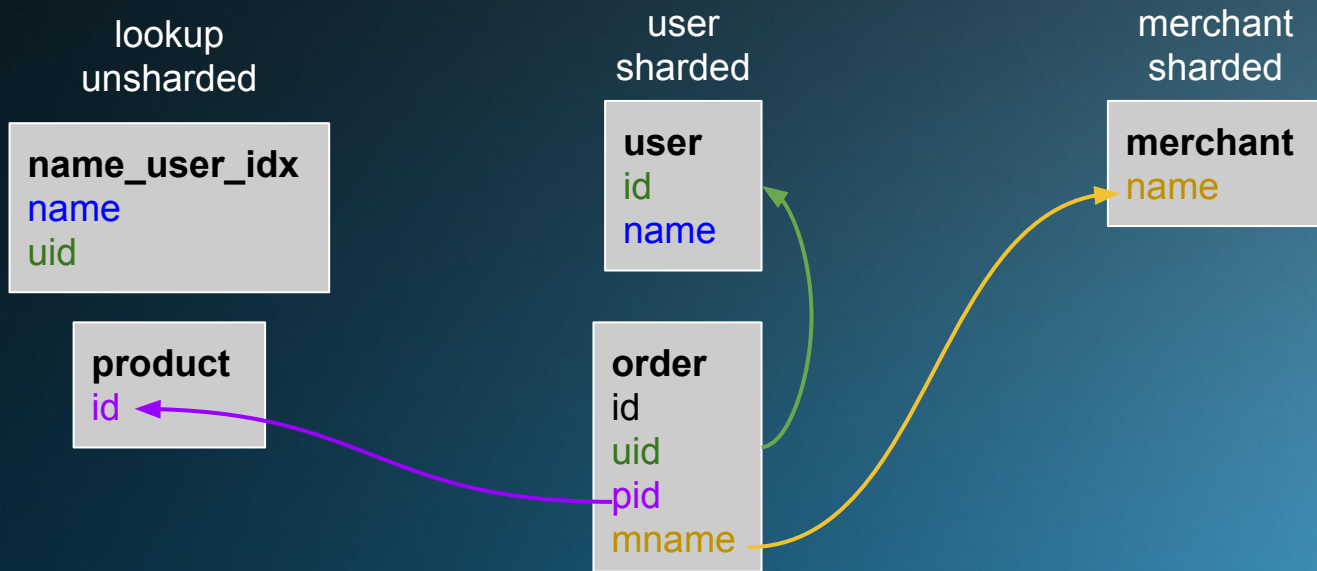


VReplication

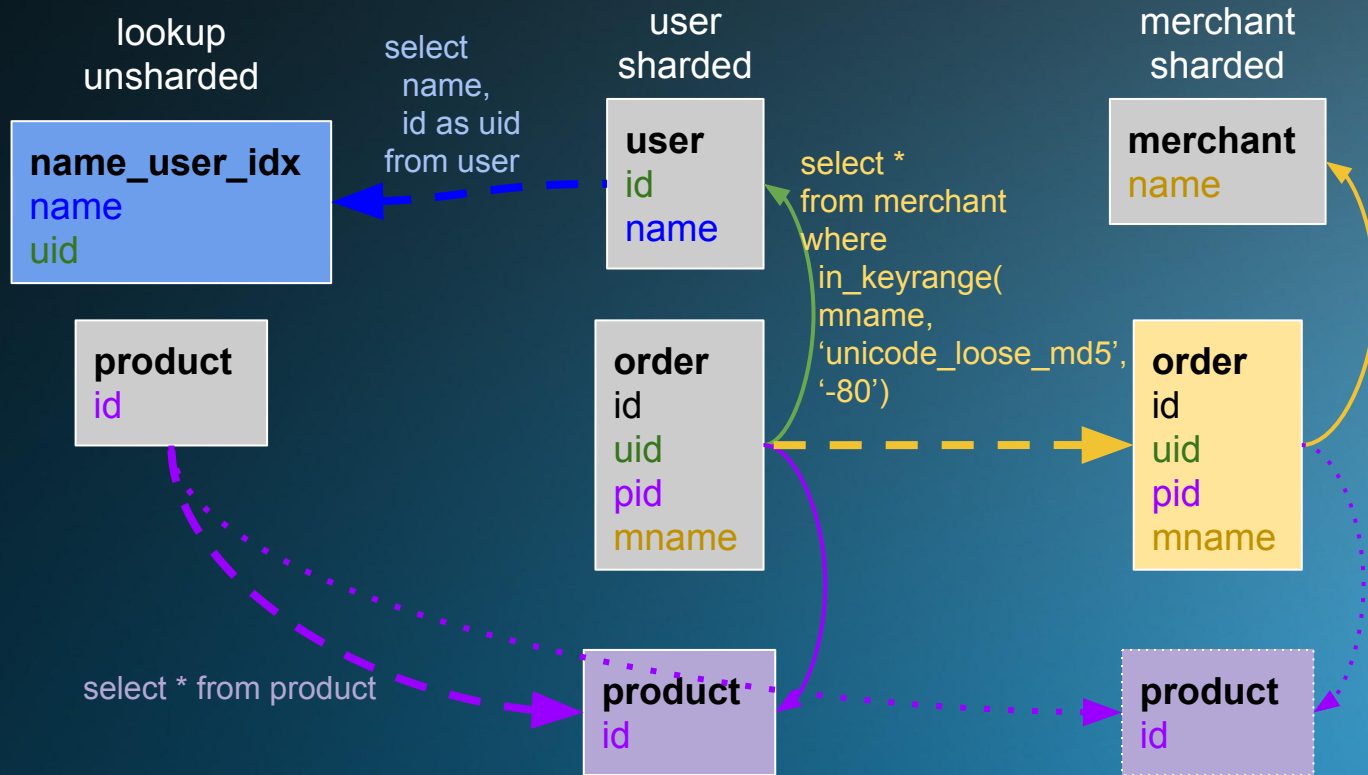
- Componentization of filtered replication
- Independent submodule
- Resharding continues to use it
- Can be used for new workflows



Demo



Demo



Roadmap

- VReplication
- Documentation
- Simplification
- Compatibility

